

ShopPilot

An Offline-First Conversational Shopping Copilot for Multi-Turn E-Commerce

Astrid CLI Engine · Zero Token Cost · Offline Default Path

TechnicalScore

0.907

Starter 0.107 · 8.5× lift

Hit Rate@10

97.5%

Starter 12.5%

MRR

0.858

Starter 0.068 · 12.6×

MTTC

2.88 turns

Starter 9.81 · -70.6%

github.com/algorathem/techjam2026-shopping-copilot

Why Keyword Search Collapses in Conversational Commerce

Single-turn search assumes static intent. Real sessions hit four core failure modes.

01

Aesthetic Under-specification

Vague queries ("something nice for summer") flood pools with flat lexical scores.

02

Semantic Collisions

Polysemy causes category drag ("dress sandals" pulling party dresses).

03

Mid-Session Preference Shifts

Pivots ("actually forget sneakers...") pollute state with stale tokens.

04

Turn Fatigue & Horizon

Abandonment compounds each unnecessary ask; sessions fail at the ≤ 10 -turn limit.

OBJECTIVE FUNCTION

$$\text{TechnicalScore} = 0.50 \cdot \text{Hit@10} + 0.30 \cdot \text{MRR} + 0.20 \cdot ((11 - \text{MTTC}) / 10)$$

ShopPilot Engine: In-Memory Multi-Turn Pipeline

Every response is a new turn in the conversation. The pipeline is designed to be stateful and context-aware.

1

Intent & Entity Triage

Regex triage: buy/browse, audience, taxonomy anchors.

2

Provenance-Aware DST

Slots separate disclosed facts from transient soft prefs.

3

In-Memory Hybrid Recall

SQLite FTS5 (BM25) fused with NumPy char n-gram dense hash.

4

Dynamic Clarification Ladder

Asymmetric other-first funnel; skip satisfied slots.

5

Constraint-Coverage Reranker

Exact-match bonuses + leaf-category + audience align.

State Is the Product: Non-Monotonic Dialog Tracking

Override hygiene · taxonomy locks · empirically validated ask policy · zero-token default.

CORE INVARIANT

Turn t₀ ("black") ranks against [dress + plus-size + black] — never "black" in isolation.

Selective Soft-Only Invalidation

On override: wipe soft prefs; keep disclosed facts the simulator will not re-send.

Taxonomy Boundary Lock

Disambiguate at ingest: dress sandals → family:footwear (not garment dresses).

Empirically Validated Ask Policy

other-first funnel beats naive max-IG (max-IG dropped Tech to ≈0.72).

Zero-Token Footprint

In-memory stdlib + NumPy. MiniLM / LLM paths are optional flags only.

Official Benchmark: 8.5× Lift Over Baseline

Deterministic evaluator.local_evaluator · Amazon CSJ 50k · N=200 · SHOPPILOT_DENSE=hash

Metric	Weak BM25	ShopPilot (hash)	Delta / Lift
TechnicalScore	0.107	0.907	+0.800 (8.5×)
Hit Rate@10	12.5%	97.5%	+85.0 pp
MRR	0.068	0.858	+0.790 (12.6×)
MTTC (turns)	9.81	2.88	-6.93 (-70.6%)
LLM Token Cost	0	0	100% offline

METRIC TRADE-OFF · THE EARLY-TERMINATION PARADOX

Emitting Top-10 too early on Turn 1 risks a low-rank hit (e.g. Rank 8 □ MRR=0.125), freezing a poor score. Because MRR carries 1.5× the weight of MTTC in TechnicalScore, ShopPilot uses Precision Gating (Top-1 limit orders on ambiguous early turns) to gather constraints and force final conversion toward Rank 1.

Performance Breakdown Across Scenarios

High conversion maintained across buying, browse, override, and boundary (N=200).

Scenario	N	Hit@10	MRR	MTTC
Targeted Buying	80	97.5%	0.833	2.45
Open Browsing	80	96.3%	0.816	2.79
Intent Override	30	96.7%	0.866	4.20
Boundary (Don't Care)	10	100.0%	0.933	3.00
Overall Public Benchmark	200	97.5%	0.858	2.88

KEY TAKEAWAY

Preference override incurs a turn penalty (4.20 MTTC) from pivot re-elicitation, but holds 96.7% Hit@10 and 0.866 MRR via clean soft-only state invalidation.

Live Execution: Astrid Multi-Turn Terminal

```
python3 cli_chat.py --dense hash · /new /state /quit
```

astrid — dense=hash

```
$ python3 cli_chat.py --dense hash
```

Astrid · quiet clarity for every aisle

You · shoes for my son

Astrid: Got it (footwear; boys). Any color...?

□ ask · other

You · Actually forget sneakers, dress shoes size 9

Astrid: Updated (footwear; size 9). Color?

□ soft wipe · disclosed kept · family=footwear

/state → size=9 asked={other}

Vague → Clarify

dress + exploring → ask other → "plus size" locks size (no re-ask).

Taxonomy Lock

"dress sandals" → family=footwear; suppresses garment dresses.

Preference Override

running shoes → dress shoes size 9; soft wipe; size persists.

Bridging Benchmark Metrics to Commercial Value

Kit scores read as merchant KPIs — findability, rank quality, cost-to-serve, infra overhead.

Hit@10 · 97.5%

Minimizes zero-result searches and bounce across sparse catalogs.

MRR · 0.858

Positions converters above the mobile fold without endless scroll.

MTTC · 2.88 turns

Cuts turn fatigue and abandonment before checkout.

Zero Token Footprint

Edge latency, deterministic path, \$0 recurring LLM API cost.

08 · CONCLUSION & REPRODUCIBILITY

ShopPilot

Deterministic, Fast, & Open Source

0.907

TechnicalScore

97.5%

Hit@10

0.858

MRR

2.88

MTTC

0

Tokens

ONE-LINE REPRODUCTION

```
export SHOPPILOT_DENSE=hash; python3 -m evaluator.local_evaluator  
github.com/algorathem/techjam2026-shopping-copilot
```

Thank you — we welcome your questions.